



Product Requirements Document: AI-Powered PR Management System

I've created a comprehensive 31-page Product Requirements Document for your AI-powered GitHub PR management application. This document is ready to send to your engineering team (Cursor) to start development.

Key Highlights of the PRD:

Strategic Positioning

Your idea targets a **real and critical problem**: 45% of open source organizations cite maintainer burnout as their #1 challenge in 2025. However, the market is crowded with AI code reviewers. This PRD positions your product to differentiate by focusing on **PR triage, workflow automation, and maintainer wellbeing** rather than just code quality review.^{[1] [2] [3]}

Core Differentiation from Competitors

Unlike GitHub Copilot, CodeRabbit, and PR-Agent which focus on code review quality, your tool will:

- **Intelligently triage and prioritize PRs** with 6-category classification system
- **Route PRs to the right reviewers** based on expertise, workload, and availability
- **Detect and prevent maintainer burnout** through proactive monitoring
- **Provide full-repository context** via RAG (Retrieval-Augmented Generation)
- **Support progressive review workflows** allowing partial PR merges

Technical Architecture Specifications

Multi-Agent System:

- Coordinator agent orchestrating 5 specialized sub-agents
- Uses Claude 3.7 Sonnet (200K context) for main analysis
- GPT-4o Mini for fast classification tasks
- Asynchronous parallel processing to reduce latency^{[4] [5]}

RAG Implementation:

- Vector database: Pinecone (managed) or Weaviate (self-hosted)^{[6] [7]}
- 768-dimension embeddings using CodeBERT

- Hybrid search (vector + full-text) for context retrieval
- Weekly re-indexing with incremental updates on push events

Technology Stack:

- Backend: Python 3.11+ with FastAPI
- Database: PostgreSQL 15+ for metadata, Redis for caching
- Infrastructure: Docker + Kubernetes for scalable deployment
- GitHub App with JWT authentication^[8] ^[9]

Feature Set Breakdown

Phase 1 MVP (Months 1-3):

- PR classification into 6 categories
- Priority scoring (0-100 based on 6 weighted factors)
- GitHub bot command interface (@PRCoPilot /trriage)
- Basic RAG context retrieval
- Daily digest emails

Phase 2 Intelligence (Months 4-6):

- Multi-agent architecture
- Enhanced full-repo RAG context
- Smart reviewer routing
- Bulk actions for maintainers

Phase 3 Wellbeing (Months 7-9):

- Burnout detection and prevention system
- Progressive review workflow (partial PR merges)
- Advanced analytics dashboard
- Cross-repository intelligence

Phase 4 Enterprise (Months 10-12):

- Custom agent training per repository
- Compliance reporting
- API for third-party integrations
- On-premise deployment packages

Business Model

Free Tier: 75 PR analyses/month (open source)
Pro Tier: \$30/user/month (unlimited)
Enterprise: Custom pricing for >100 users

Target: 1,000 active repositories by month 6, \$50K+ MRR by month 12

Security & Privacy

- **Zero data retention** commitment (code analyzed in-memory only)
- Self-hosted option available for strict data residency requirements
- SOC 2, GDPR, CCPA compliant
- Users provide their own LLM API keys (OpenAI, Anthropic)

Risk Mitigation

The document addresses 7 major risks including:

- LLM hallucinations → Confidence thresholds + human validation
- Competition from GitHub Copilot → Focus on differentiated triage features
- Maintainer skepticism of AI → Transparent reasoning + "co-pilot" positioning

What Makes This Different from Existing Tools

The competitive analysis shows that while GitHub Copilot (\$39/mo), CodeRabbit (\$12-24/mo), and PR-Agent (free) focus heavily on **code review quality**, your tool is the **only one focused on maintainer workflow optimization**:

- ✓ **PR Triage & Prioritization** (competitors: ✕)
- ✓ **Burnout Prevention** (competitors: ✕)
- ✓ **Smart Reviewer Routing** (competitors: ✕)
- ✓ **Multi-Agent Architecture** (competitors: ✕)
- ✓ **Full-Repo RAG Context** (competitors: PR-diff only)

Ready for Development

The PRD includes everything Cursor needs to start building:

- Database schemas (PostgreSQL tables)
- API endpoint specifications
- GitHub App permissions manifest
- Prompt engineering templates
- Infrastructure deployment guides (Docker, Kubernetes)

- Complete tech stack recommendations with specific versions

The document is structured for iterative development with clear success metrics for each phase and specific milestones to validate product-market fit before scaling.

Your next step: Share this PRD with your engineering team and schedule a kickoff meeting to align on Phase 1 MVP scope and timeline. The document recommends starting with 100 beta testers from mid-size open source projects (1K-10K stars) in Python/JavaScript ecosystems.

✱✱

1. <https://www.openlogic.com/system/files/2025-05/report-openlogic-2025-state-of-open-source-support.pdf>
2. <https://opensource.org/blog/key-insights-from-the-2025-state-of-open-source-report>
3. <https://www.linkedin.com/pulse/maintainer-burnout-real-why-open-source-stagnating-peter-smulovics-9bgze>
4. <https://cloud.google.com/architecture/multiagent-ai-system>
5. <https://www.anthropic.com/engineering/built-multi-agent-research-system>
6. <https://aloha.co/ai/comparisons/vector-database-comparison/pinecone-vs-weaviate>
7. <https://xenoss.io/blog/vector-database-comparison-pinecone-qdrant-weaviate>
8. <https://docs.github.com/en/apps/creating-github-apps/registering-a-github-app/using-webhooks-with-github-apps>
9. <https://notes.kodekloud.com/docs/AZ-400/Design-and-Implement-Authentication-and-Authorization-Methods/Implement-and-manage-GitHub-Authentication>
10. <https://docs.github.com/en/rest/apps/webhooks>
11. <https://www.keyreply.com/blog/conversational-design-virtual-assistants>
12. <https://dl.acm.org/doi/10.1145/3663533.3664041>
13. <https://www.anaconda.com/blog/my-open-source-journey-from-graduation-to-maintainership>
14. <https://determ.com/blog/pr-automation-best-tools-which-do-the-work-for-you/>
15. <https://docs.github.com/en/organizations/keeping-your-organization-secure/managing-two-factor-authentication-for-your-organization/managing-bots-and-service-accounts-with-two-factor-authentication>
16. <https://www.meltwater.com/en/blog/what-does-pr-automation-mean-for-pr-pros>
17. https://docs.github.com/en/actions/tutorials/authenticate-with-github_token
18. <https://docs.github.com/en/webhooks/webhook-events-and-payloads>
19. <https://thechigroup.co/articles/2025/1/10/pr-automation-explained-how-to-save-time-and-maximize-your-pr-impact>
20. <https://docs.github.com/actions/security-guides/automatic-token-authentication>
21. <https://github.blog/changelog/2025-07-01-enterprise-level-access-for-github-apps-and-installation-automation-apis/>
22. <https://www.agilitypr.com/pr-news/pr-skills-profession/how-automation-helps-pr-teams-maintain-consistent-messaging-across-channels/>
23. <https://docs.github.com/en/rest/apps>
24. <https://www.axiapr.com/blog/5-ways-to-automate-pr-processes-for-streamlined-lead-generation>

25. <https://github.com/orgs/community/discussions/67874>
26. <https://docs.github.com/en/apps/creating-github-apps/about-creating-github-apps/about-creating-github-apps>
27. <https://prprofessionals.in/blog/how-pr-automation-tools-maximize-client-engagement>
28. <https://docs.github.com/en/code-security/supply-chain-security/end-to-end-supply-chain/securing-accounts>
29. <https://docs.github.com/en/rest>
30. <https://wprd.app/the-role-of-ai-and-automation-in-public-relations-pr-hr-implications-and-opportunities-by-adedeji-adeniyi/>
31. <https://marutitech.com/how-to-choose-right-llm/>
32. <https://nexla.com/ai-infrastructure/retrieval-augmented-generation/>
33. <https://himalayas.app/companies/github/tech-stack>
34. <https://huggingface.co/blog/dvilasuero/choosing-best-open-source-ai-models>
35. <https://www.merge.dev/blog/rag-best-practices>
36. <https://github.com/stackshareio/awesome-stacks>
37. <https://www.linkgraph.com/blog/best-llm-model/>
38. <https://learn.microsoft.com/en-us/azure/architecture/ai-ml/guide/rag/rag-solution-design-and-evaluation-guide>
39. <https://www.imaginarycloud.com/blog/techstack-mobile-app>
40. <https://arxiv.org/html/2508.00083v1>
41. <https://orkes.io/blog/rag-best-practices/>
42. <https://github.com/orgs/community/discussions/167234>
43. <https://www.flowhunt.io/blog/best-llms-for-coding-june-2025/>
44. <https://python.langchain.com/docs/tutorials/rag/>
45. <https://syndicode.com/blog/how-to-choose-tech-stack/>
46. <https://www.instaclustr.com/education/open-source-ai/top-10-open-source-llms-for-2025/>
47. <https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval>
48. <https://github.com/topics/tech-stack>
49. <https://dextralabs.com/blog/best-llm-for-coding/>
50. <https://aws.amazon.com/what-is/retrieval-augmented-generation/>
51. <https://aws.amazon.com/blogs/modernizing-with-aws/automate-microsoft-web-application-deployments-with-github-actions-and-terraform/>
52. <https://www.koyeb.com/docs/build-and-deploy/deploy-with-git>
53. <https://www.cognizant.com/us/en/ai-lab/blog/multi-agent-evaluation-system>
54. <https://www.howtobuysaas.com/blog/pinecone-vs-weaviate-vs-chroma/>
55. <https://docs.github.com/actions/deployment/about-deployments/deploying-with-github-actions>
56. <https://research.aimultiple.com/vector-database-for-rag/>
57. <https://learn.microsoft.com/en-us/azure/architecture/ai-ml/guide/ai-agent-design-patterns>
58. <https://lakefs.io/blog/12-vector-databases-2023/>
59. <https://learn.microsoft.com/en-us/azure/app-service/deploy-github-actions>

60. <https://www.docker.com/blog/how-to-build-a-multi-agent-system/>
61. <https://liveblocks.io/blog/whats-the-best-vector-database-for-building-ai-products>
62. <https://docs.github.com/en/apps/creating-github-apps/about-creating-github-apps/deciding-when-to-build-a-github-app>
63. <https://www.nitorinfotech.com/blog/multi-agent-collaboration-how-ai-agents-work-together/>
64. <https://dev.to/dandv/how-to-choose-a-vector-database-pinecone-weaviate-mongodb-atlas-semanticdb-a09>
65. <https://codefresh.io/learn/github-actions/deployment-with-github-actions/>
66. <https://productschool.com/blog/product-strategy/product-template-requirements-document-prd>
67. <https://docs.github.com/en/apps/creating-github-apps/about-creating-github-apps/best-practices-for-creating-a-github-app>
68. <https://www.qodo.ai/blog/automated-code-review/>
69. <https://www.atlassian.com/agile/product-management/requirements>
70. <https://github.com/security/advanced-security>
71. <https://coderabbit.ai>
72. <https://www.notion.com/templates/category/product-requirements-doc>
73. <https://www.checkpoint.com/cyber-hub/cloud-security/what-is-developer-security/21-security-best-practices-for-github/>
74. <https://research.aimultiple.com/ai-code-review-tools/>
75. <https://www.aha.io/roadmapping/guide/requirements-management/what-is-a-good-product-requirements-document-template>
76. <https://github.blog/changelog/2025-06-24-security-updates-for-apps-and-api-access/>
77. <https://github.com/resources/articles/ai/ai-code-reviews>
78. <https://www.figma.com/resource-library/product-requirements-document/>
79. <https://github.blog/changelog/2025-03-04-introducing-github-secret-protection-and-github-code-security/>
80. <https://getdx.com/blog/ai-coding-assistant-pricing/>
81. <https://www.atlassian.com/software/confluence/templates/product-requirements>
82. <https://www.gitguardian.com/glossary/what-are-github-security-best-practices>
83. <https://www.greptile.com>
84. <https://www.hustlebadger.com/what-do-product-teams-do/prd-template-examples/>
85. <https://docs.github.com/en/apps/github-marketplace/creating-apps-for-github-marketplace/security-best-practices-for-apps-on-github-marketplace>
86. <https://docs.github.com/en/code-security/dependabot/dependabot-auto-triage-rules/customizing-auto-triage-rules-to-prioritize-dependabot-alerts>
87. https://www.reddit.com/r/opensource/comments/1ciq1xo/open_source_maintainers_tell_me_about_your/
88. <https://www.conversationdesigninstitute.com/topics/best-practices>
89. <https://docs.github.com/en/code-security/securing-your-organization/understanding-your-organizations-exposure-to-vulnerabilities/prioritizing-dependabot-alerts-using-metrics>
90. <https://www.jlowin.dev/blog/oss-maintainers-guide-to-saying-no>
91. <https://botpress.com/blog/conversation-design>

92. <https://docs.github.com/en/code-security/code-scanning/managing-code-scanning-alerts/triaging-code-scanning-alerts-in-pull-requests>
93. <https://nolanlawson.com/2017/03/05/what-it-feels-like-to-be-an-open-source-maintainer/>
94. <https://developers.google.com/assistant/conversation-design/learn-about-conversation>
95. <https://james-vu.com/blog/f/step-by-step-guide-to-triage-security-findings-in-github>
96. <https://antirez.com/news/129>
97. <https://www.salesforce.com/blog/what-is-conversation-design/>
98. <https://docs.github.com/en/code-security/dependabot/dependabot-auto-triage-rules>
99. <https://opensauced.pizza/blog/when-open-source-maintainers-leave>
100. <https://rasa.com/blog/how-to-design-chatbot-conversation/>
101. <https://docs.github.com/en/code-security/security-overview/viewing-metrics-for-pull-request-alerts>
102. <https://openssf.org/blog/2024/01/31/maintainer-motivations-challenges-and-best-practices-on-open-source-software-security/>